

```

# Time Series Prediction using Deep Learning With LSTM
# -----

import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.model_selection import train_test_split
from keras.preprocessing.sequence import TimeseriesGenerator
import tensorflow as tf

mpl.rcParams['figure.figsize'] = (10,8)
mpl.rcParams['axes.grid'] = False

df = pd.read_csv('/retail_sales_Time_series.csv')
df

{"summary": "{\n  \"name\": \"df\",\n  \"rows\": 52585,\n  \"fields\": [\n    {\n      \"column\": \"Date\",\n      \"properties\": {\n        \"dtype\": \"object\",\n        \"num_unique_values\": 52585,\n        \"samples\": [\n          \"30/4/18 10:00\",\n          \"12/11/17 12:00\",\n          \"16/6/16 22:00\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Sales\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 29.957814641460093,\n        \"min\": 32.79007417,\n        \"max\": 168.1072812,\n        \"num_unique_values\": 52585,\n        \"samples\": [\n          105.6581726,\n          96.49120044,\n          72.44353151\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"df\"}

df.info()


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 52585 entries, 0 to 52584
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Date    52585 non-null    datetime64[ns]
1    Sales    52585 non-null    float64
dtypes: datetime64[ns](1), float64(1)
memory usage: 821.8 KB

df['Date'] = pd.to_datetime(df['Date'], infer_datetime_format=True)

<ipython-input-8-45a55d54a14d>:1: UserWarning: The argument
'infer_datetime_format' is deprecated and will be removed in a future
version. A strict version of it is now the default, see
https://pandas.pydata.org/pdeps/0004-consistent-to-datetime-

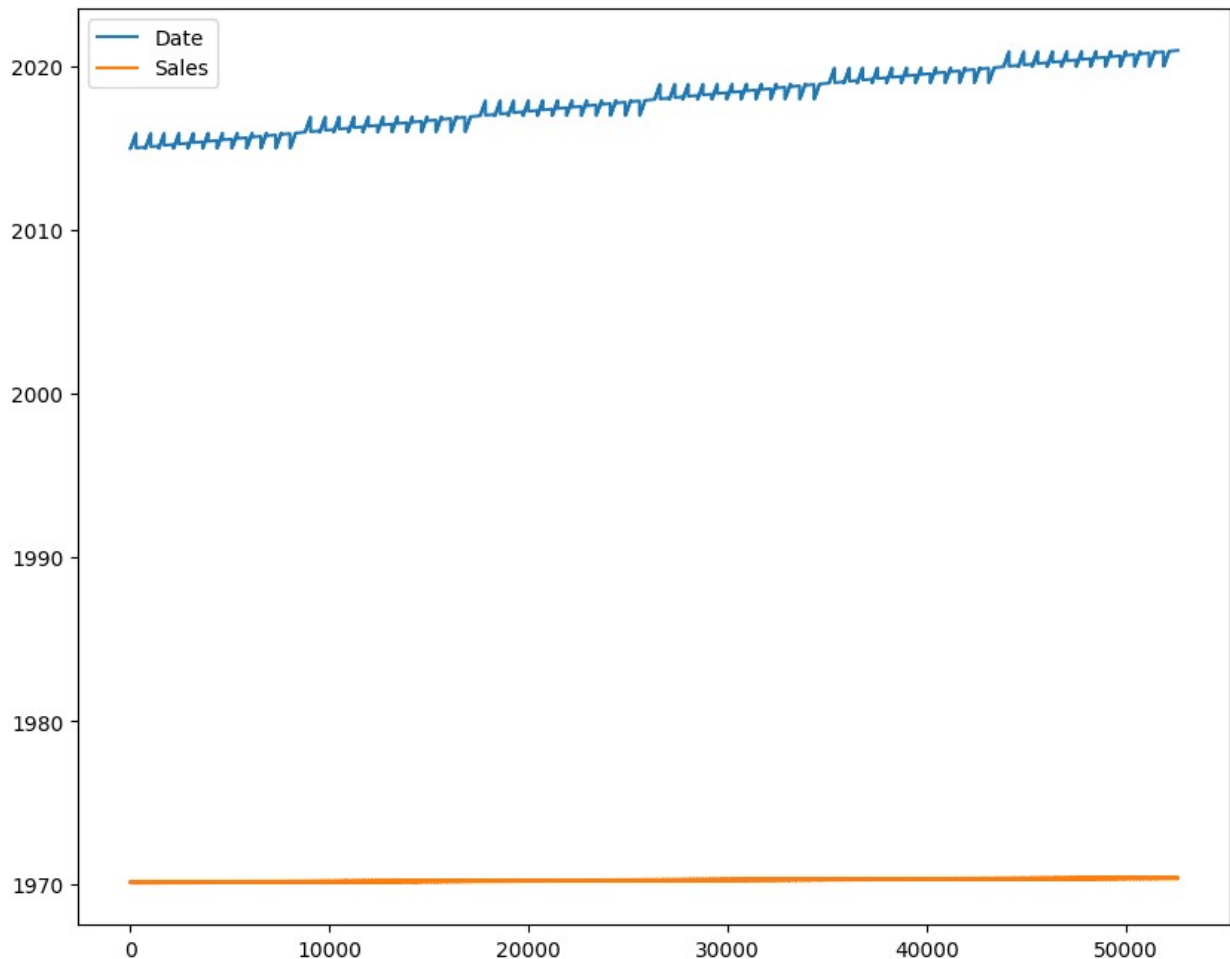
```

```
parsing.html. You can safely remove this argument.  
df['Date'] = pd.to_datetime(df['Date'], infer_datetime_format=True)  
<ipython-input-8-45a55d54a14d>:1: UserWarning: Could not infer format,  
so each element will be parsed individually, falling back to  
'dateutil'. To ensure parsing is consistent and as-expected, please  
specify a format.  
df['Date'] = pd.to_datetime(df['Date'], infer_datetime_format=True)  
df.set_index('Date')[['Sales']].plot(subplots=True)  
array([<Axes: xlabel='Date'>], dtype=object)
```



```
df.plot(figsize=(10,8))
```

```
<Axes: >
```



```
df_input=df[['Sales']]
df_input
```

```
{
  "summary": {
    "name": "df_input",
    "rows": 52585,
    "fields": [
      {
        "column": "Sales",
        "properties": {
          "dtype": "number",
          "std": 29.957814641460093,
          "min": 32.79007417,
          "max": 168.1072812,
          "num_unique_values": 52585
        }
      }
    ],
    "samples": [
      [
        105.6581726,
        96.49120044,
        72.44353151
      ]
    ],
    "semantic_type": "",
    "description": ""
  },
  "type": "dataframe",
  "variable_name": "df_input"
}
```

```
df_input.describe()
```

```
{
  "summary": {
    "name": "df_input",
    "rows": 8,
    "fields": [
      {
        "column": "Sales",
        "properties": {
          "dtype": "number",
          "std": 18559.801541552788,
          "min": 29.957814641460093,
          "max": 52585.0,
          "num_unique_values": 8
        }
      }
    ],
    "samples": [
      [
        99.9974882959477,
        99.93064275
      ]
    ]
  }
}
```

```
52585.0\n                ],\n                \n                \"semantic_type\": \"\",\n                }\n            }\n        ]\n    }\", \"type\": \"dataframe\"}\n\n    df_input.query('Sales>124')\n\n    {\"summary\": \"{\\n    \"name\": \"df_input\\n\",\\n    \"rows\": 13620,\\n    \"fields\": [\\n        {\\n            \"column\": \"Sales\\n\",\\n            \"properties\": {\\n                \"dtype\": \"number\\n\",\\n                \"std\": 9.308234762176067,\\n                \"min\": 124.0014659,\\n                \"max\": 168.1072812,\\n                \"num_unique_values\": 13620,\\n                \"samples\": [\\n                    128.09031,\\n                    132.9555999,\\n                    148.970802\\n                ],\\n                \"semantic_type\": \"\",\\n                }\n            }\n        ]\n    }\", \"type\": \"dataframe\"}\n\n    scaler =MinMaxScaler()\n    data_scaled = scaler.fit_transform(df_input)\n\n    data_scaled98\n\n    array([[0.15710157],\n           [0.1283763 ],\n           [0.17118935],\n           ...,\n           [0.82542345],\n           [0.85059557],\n           [0.86783799]])\n\n    features=data_scaled\n    target=data_scaled[:,0]\n\n    TimeseriesGenerator(features,target,\n    length=2,sampling_rate=1,batch_size=1)[0]\n\n    (array([[0.15710157],\n           [0.1283763 ]]),\n      array([0.17118935]))\n\n    x_train, x_test, y_train, y_test = train_test_split(features ,target,\n    test_size=0.20, random_state=125, shuffle=False)\n\n    x_train.shape\n\n    (42068, 1)\n\n    X_test.shape\n\n    (10517, 1)\n\n    win_length = 250\n    batch_size = 32\n    num_features = 1\n    train_generator = TimeseriesGenerator(x_train, y_train,
```

```
length=win_length, sampling_rate=1, batch_size=batch_size)
test_generator = TimeseriesGenerator(x_test , y_test,
length=win_length, sampling_rate=1, batch_size=batch_size)
```

```
train_generator[0]
```

```
(array([[0.15710157],
        [0.1283763 ],
        [0.17118935],
        ...,
        [0.18276762],
        [0.22175255],
        [0.20120757]]),

      [[0.1283763 ],
        [0.17118935],
        [0.17561098],
        ...,
        [0.22175255],
        [0.20120757],
        [0.14312201]]),

      [[0.17118935],
        [0.17561098],
        [0.20468793],
        ...,
        [0.20120757],
        [0.14312201],
        [0.15166137]]),

      ...,

      [[0.22315337],
        [0.24572904],
        [0.22854168],
        ...,
        [0.10166027],
        [0.06828254],
        [0.06927693]]),

      [[0.24572904],
        [0.22854168],
        [0.2129003 ],
        ...,
        [0.06828254],
        [0.06927693],
        [0.0739013 ]],

      [[0.22854168],
        [0.2129003 ],
```

```

        [0.22877846],
        ...,
        [0.06927693],
        [0.0739013 ],
        [0.01741361]])),
array([0.14312201, 0.15166137, 0.1125318 , 0.10310585, 0.08499234,
       0.06207513, 0.04333615, 0.03926781, 0.03596528, 0.04262377,
       0.0642177 , 0.06894608, 0.10678641, 0.13736663, 0.10557498,
       0.13641166, 0.15753036, 0.17157231, 0.176193 , 0.19168563,
       0.18814053, 0.15758967, 0.17846883, 0.16403583, 0.12151806,
       0.11309371, 0.10166027, 0.06828254, 0.06927693, 0.0739013 ,
       0.01741361, 0.00389642]))

```

```

model =tf.keras.Sequential()
model.add(tf.keras.layers.LSTM(128, input_shape =
(win_length,num_features), return_sequences=True))
model.add(tf.keras.layers.LeakyReLU(alpha=0.5))
model.add(tf.keras.layers.LSTM(128,return_sequences=True))
model.add(tf.keras.layers.LeakyReLU(alpha=0.5))
model.add(tf.keras.layers.Dropout(0.3))
model.add(tf.keras.layers.LSTM(64,return_sequences=False))
model.add(tf.keras.layers.LeakyReLU(alpha=0.3))
model.add(tf.keras.layers.Dense(1))

```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
lstm (LSTM)	(None, 250, 128)	66560
leaky_re_lu (LeakyReLU)	(None, 250, 128)	0
lstm_1 (LSTM)	(None, 250, 128)	131584
leaky_re_lu_1 (LeakyReLU)	(None, 250, 128)	0
dropout (Dropout)	(None, 250, 128)	0
lstm_2 (LSTM)	(None, 64)	49408
leaky_re_lu_2 (LeakyReLU)	(None, 64)	0
dense (Dense)	(None, 1)	65

```
=====
Total params: 247617 (967.25 KB)
```

```
Trainable params: 247617 (967.25 KB)
```

Non-trainable params: 0 (0.00 Byte)

```
early_stopping =  
tf.keras.callbacks.EarlyStopping(monitor='val_loss',patience=2,mode =  
'min')
```

```
model.compile(loss= tf.losses.MeanSquaredError(),optimizer =  
tf.optimizers.Adam(),metrics=[tf.metrics.MeanAbsoluteError()])
```

```
history = model.fit_generator(train_generator, epochs=50,  
validation_data=test_generator, shuffle=False,  
callbacks=[early_stopping])
```

Epoch 1/50

```
<ipython-input-48-85f216b38176>:5: UserWarning: `Model.fit_generator`  
is deprecated and will be removed in a future version. Please use  
`Model.fit`, which supports generators.
```

```
history = model.fit_generator(train_generator, epochs=50,  
validation_data=test_generator, shuffle=False,  
callbacks=[early_stopping])
```

```
1307/1307 [=====] - 1215s 925ms/step - loss:  
8.7724e-04 - mean_absolute_error: 0.0223 - val_loss: 8.7840e-04 -  
val_mean_absolute_error: 0.0232
```

Epoch 2/50

```
1307/1307 [=====] - 1514s 1s/step - loss:  
0.0013 - mean_absolute_error: 0.0184 - val_loss: 0.0022 -  
val_mean_absolute_error: 0.0385
```

Epoch 3/50

```
1307/1307 [=====] - 1343s 1s/step - loss:  
3.9235e-04 - mean_absolute_error: 0.0154 - val_loss: 9.4191e-04 -  
val_mean_absolute_error: 0.0246
```

```
model.evaluate_generator(test_generator, verbose=0)
```

```
<ipython-input-49-a68ec1088d46>:1: UserWarning:  
`Model.evaluate_generator` is deprecated and will be removed in a  
future version. Please use `Model.evaluate`, which supports  
generators.
```

```
model.evaluate_generator(test_generator, verbose=0)
```

```
[0.0009419145644642413, 0.024641329422593117]
```

```
predictions=model.predict_generator(test_generator)
```

```
<ipython-input-50-6ca49477046f>:1: UserWarning:  
`Model.predict_generator` is deprecated and will be removed in a  
future version. Please use `Model.predict`, which supports generators.  
predictions=model.predict_generator(test_generator)
```

```
predictions.shape[0]
```

```
10267
```

```
predictions
```

```
array([[0.77706367],
       [0.7703783 ],
       [0.75869006],
       ...,
       [0.7953903 ],
       [0.8085752 ],
       [0.82290584]], dtype=float32)
```

```
y_test
```

```
array([0.694982 , 0.67728269, 0.71390419, ..., 0.82542345,
       0.85059557,
       0.86783799])
```

```
x_test
```

```
array([[0.694982 ],
       [0.67728269],
       [0.71390419],
       ...,
       [0.82542345],
       [0.85059557],
       [0.86783799]])
```

```
x_test[:,0:][win_length:]
```

```
array([[0.76973 ],
       [0.75606292],
       [0.76404814],
       ...,
       [0.82542345],
       [0.85059557],
       [0.86783799]])
```

```
df_pred=pd.concat([pd.DataFrame(predictions),pd.DataFrame(x_test[:,1:]
[win_length:])],axis=1)
```

```
df_pred
```

```
{"summary":{"name": "df_pred", "rows": 10267,
"fields": [{"column": 0, "properties": {"dtype": "float32", "num_unique_values": 10257,
"samples": [0.6808833479881287,
0.7794239521026611,
0.8592292666435242],
"semantic_type": "", "description": ""}
}]},"type":"dataframe","variable_name":"df_pred"}
```



```

rev_trans=scaler.inverse_transform(df_pred)
rev_trans
array([[137.94017],
       [137.03552],
       [135.4539 ],
       ...,
       [140.42007],
       [142.20422],
       [144.1434 ]], dtype=float32)

df_final=df_input[predictions.shape[0]*-1:]
df_final.count()
Sales      10267
dtype: int64

df_final['Sales_pred']=rev_trans[:,0]

<ipython-input-66-1d01bceaffbb>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
df_final['Sales_pred']=rev_trans[:,0]

df_final

{"summary":{"name": "df_final", "rows": 10267, "fields": [{"column": "Sales", "properties": {"dtype": "number", "std": 9.809966216969036, "min": 111.7140685, "max": 168.1072812, "num_unique_values": 10267, "samples": [{"value": 158.4069433, "count": 155.9339592, "semantic_type": "", "description": ""}], "column": "Sales_pred", "properties": {"dtype": "float32", "num_unique_values": 10247, "samples": [{"value": 138.82313537597656, "count": 123.15750885009766, "semantic_type": "", "description": ""}], "column": ""}], "type": "dataframe", "variable_name": "df_final"}

df_final[['Sales', 'Sales_pred']].plot()

<Axes: >

```

